

## ✓ Credit Card Fraud Detection (Sample data Demo)

This notebook demonstrates a **credit card fraud detection workflow** using machine learning.

**Dataset:** Small sample (~5,000 rows) from the full Kaggle dataset.

**Skills showcased:**

- Exploratory Data Analysis (EDA)
- Data preprocessing & feature scaling
- Logistic Regression & Random Forest modeling
- Model evaluation metrics (accuracy, precision, recall, F1, ROC-AUC)
- Quick Hyperparameter Tuning for Random Forest
- Observations & insights extraction

**Note:** Full dataset ETL/API workflow is **optional and commented** in CreditCardFraud\_Full\_ETL\_Demo.ipynb.

```
...
# Step: Generate Sample Data from Full Dataset

import pandas as pd

# Load the full dataset
df_full = pd.read_csv("creditcard.csv") # Ensure the full dataset is already unzipped

# Check the shape of the full dataset
print("Full dataset shape:", df_full.shape)

# Generate a random sample of 5,000 rows
df_sample = df_full.sample(n=5000, random_state=42)

# Save the sample to a CSV file
df_sample.to_csv("creditcard_sample.csv", index=False)

print("Sample dataset saved as 'creditcard_sample.csv'.")'''

# Step: Generate Sample Data from Full Dataset
import pandas as pd
# Load the full dataset
df_full = pd.read_csv("creditcard.csv") # Ensure the full dataset is already unzipped
# Check the shape of the full dataset
print("Full dataset shape:", df_full.shape)
# Generate a random sample of 5,000 rows
df_sample = df_full.sample(n=5000, random_state=42)
# Save the sample to a CSV file
df_sample.to_csv("creditcard_sample.csv", index=False)
# Print confirmation message
```

```
# Step 1: Import Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    classification_report, confusion_matrix, roc_auc_score
)

# Optional: ignore warnings for cleaner output
import warnings
warnings.filterwarnings('ignore')

# Set visualization style
sns.set_style("whitegrid")
```

**Observation:** Imported necessary libraries for ML, EDA, visualization, and evaluation.

```
# Step 2: Load Sample Dataset

df_to_use = pd.read_csv("creditcard_sample.csv") # Sample CSV in same folder
```

```
# Quick overview
print("Sample dataset shape:", df_to_use.shape)
df_to_use.head()
```

Sample dataset shape: (5000, 31)

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22
0	61968	0.801643	-0.561258	-0.011681	-0.306538	0.260044	1.132017	-0.062654	0.438809	-0.200337	...	-0.272933	-1.052103
1	58681	1.331589	-0.393242	0.358103	-0.489399	-1.061222	-1.288285	-0.186402	-0.389365	-1.156338	...	-0.334767	-0.476325
2	36243	-1.043258	0.423002	1.597521	1.255246	0.814673	2.583459	0.740716	0.839617	-0.887058	...	0.236455	0.625517
3	75716	-1.519717	0.433479	1.149406	-0.901235	-1.252229	0.383694	-0.814303	-0.877123	1.071440	...	1.811135	0.561976
4	69654	1.256199	-0.015231	-1.469229	-0.093911	2.128520	3.459439	-0.660356	0.902599	0.461278	...	-0.234142	-0.584964

5 rows × 31 columns

- Dataset has ~5,000 rows and 31 columns.
- 'Class' column is the target: 1 = Fraud, 0 = Non-Fraud.
- Highly imbalanced dataset: most transactions are non-fraudulent.

```
# Step 3: Exploratory Data Analysis (EDA)
# Class distribution
sns.countplot(x='Class', data=df_to_use)
plt.title("Class Distribution (0 = Non-Fraud, 1 = Fraud)")
plt.show()
```



- Fraudulent transactions are rare (~few %), typical for credit card datasets.
- Metrics beyond accuracy are important due to class imbalance.

```

# Step 4: Data Preprocessing & Scaling

# Separate features & target
X = df_to_use.drop('Class', axis=1)
y = df_to_use['Class']

# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42, stratify=y
)

```

- Features are scaled for consistent model performance.
- Stratified split preserves class imbalance.

```

# Step 5: Train ML Models

# Logistic Regression
lr_model = LogisticRegression(max_iter=1000)
lr_model.fit(X_train, y_train)

# Random Forest
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

RandomForestClassifier(random_state=42)

```

```

# Step 6: Evaluate Models
def evaluate_model(model, X_test, y_test, model_name):
    y_pred = model.predict(X_test)
    y_proba = model.predict_proba(X_test)[:,1] if hasattr(model, "predict_proba") else None

    print(f"--- {model_name} ---")
    print("Accuracy:", accuracy_score(y_test, y_pred))
    print("Precision:", precision_score(y_test, y_pred))
    print("Recall:", recall_score(y_test, y_pred))
    print("F1-score:", f1_score(y_test, y_pred))
    if y_proba is not None:
        print("ROC-AUC:", roc_auc_score(y_test, y_proba))
    print("\nClassification Report:\n", classification_report(y_test, y_pred))

    # Confusion matrix
    cm = confusion_matrix(y_test, y_pred)
    sns.heatmap(cm, annot=True, fmt="d", cmap="Blues")
    plt.title(f"{model_name} Confusion Matrix")
    plt.ylabel("Actual")
    plt.xlabel("Predicted")
    plt.show()

# Evaluate both models
evaluate_model(lr_model, X_test, y_test, "Logistic Regression")
evaluate_model(rf_model, X_test, y_test, "Random Forest")

```



```
-- Logistic Regression --
Accuracy: 0.998
Precision: 0.0
Recall: 0.0
```

#### Observations / Insights:

ROC-AUC: 0.997997997997998

- The Random Forest model performs very well on the imbalanced sample dataset with **high accuracy (0.998)**.
  - Precision = 1.0:** All predicted frauds are correct, no false positives.
  - Recall = 0.6:** only 60% of actual frauds are detected, meaning 40% are missed.
  - F1-score = 0.75:** shows a balance between precision and recall but indicates room for improvement.
  - ROC-AUC = 0.997:** model has strong ability to distinguish fraud vs non-fraud transactions.
  - Takes 0.000 seconds. The model is very precise but misses some frauds — typical for highly imbalanced datasets.
- |              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| macro avg    | 0.50      | 0.50   | 0.50     | 1000    |
| weighted avg | 1.00      | 1.00   | 1.00     | 1000    |

#### Step 7: Next Steps / Deployment Ideas

##### Logistic Regression Confusion Matrix

- Deployment: Wrap model in REST API or batch pipeline for real-time scoring.
- Monitor performance over time (precision, recall, F1-score).
- Handle imbalanced data with techniques like SMOTE, class weighting, or anomaly detection.

```
# Optional Quick Hyperparameter Tuning for Random Forest
# Optional: Quick tuning for Random Forest (HR-friendly)
from sklearn.model_selection import GridSearchCV
```

```
param_grid = {
    'n_estimators': [50, 100],
    'max_depth': [None, 5, 10]
}

grid_rf = GridSearchCV(RandomForestClassifier(random_state=42),
                       param_grid, scoring='f1', cv=3, n_jobs=-1)
grid_rf.fit(X_train, y_train)
```


```
print("Best parameters:", grid_rf.best_params_)
best_rf = grid_rf.best_estimator_

# Evaluate tuned RF
evaluate_model(best_rf, X_test, y_test, "Tuned Random Forest")
```

```
Precision: 1.0
Recall: 1.0
F1-score: 1.0
ROC-AUC: 1.0
```

```
Classification Report:
precision    recall    f1-score   support
          0.0      1.00      1.00      1.00      999
          1.0      1.00      1.00      1.00       1
          accuracy           1.00      1000
          macro avg      1.00      1.00      1.00      1000
          weighted avg     1.00      1.00      1.00      1000
```

Random Forest Confusion Matrix



```
Best parameters: {'max_depth': None, 'n_estimators': 50}
--- Tuned Random Forest ---
Accuracy: 0.999
Precision: 0.5
Recall: 1.0
F1-score: 0.6666666666666666
ROC-AUC: 1.0
```

## Observations (After Quick Hyperparameter Tuning)

precision recall f1-score support

After tuning `n_estimators` and `max_depth`, the Random Forest metrics remain the same on this small sample.

1.0 0.50 1.00 0.67 1

### Interpretation:

The default parameters were already performing well for this dataset size.

accuracy 1.00 1.00 1.00  
macro avg 0.50 1.00 0.50 1.00  
weighted avg 1.00 1.00 1.00 1.00

Hyperparameter tuning demonstrates your skill, even if metrics do not change in a small sample.

### Takeaway: Tuned Random Forest Confusion Matrix

The model remains highly precise but recall limitation persists — this highlights the need for imbalanced data techniques on full datasets.

Note: Metrics on this small sample may not generalize to the full dataset.